Launch an EC2 Instance

Hadoop is built for running on commodity hardware in a data center. However, I imagine few of you have server racks in your garage. Since you don't have access to a bunch of hardware, you're going to deploy a Hadoop cluster on Amazon Web Service (AWS) EC2 (Elastic Compute Cloud) virtual machines.

These are virtual machines we can request and run from anywhere in the world, typically at some cost per hour, per machine. Each virtual machine, an instance for shorthand, has a configurable number of CPUs and amount of memory and storage. We're going to be using a free offering, but most organizations will need more resources to store and analyze their data.

If you don't have an AWS account, you'll need to create one. It might ask you for your credit card information. This is because the instances are charged by the hour. For this first lesson, you'll be using free EC2 instances. However, completing the next two lessons will require using non-free instances.

After you log in, you should choose a region that is closest to you shown here.



Now we're ready to create some instances. Here's what you should do:

- 1. From "Services", select "EC2".
- 2. Click the "Launch Instance" button.
- 3. Choose Ubuntu 64-Bit.
- 4. Choose t2.micro. Click "Next: Configure Instance Details."
- 5. Leave the number of instances to 1. We're going to install Hadoop on this instance and use it as a generic Hadoop node. Click "Next: Add Storage"

- 6. This is where you set the storage for your instance, 8 GB is fine. Of course, if you're doing this for an organization, you'll need to choose an appropriate storage size to fit your data. Click "Next: Tag Instance"
- 7. Here you can assign a tag to your instance. Anything is fine, I used "Hadoop node". Click "Next: Configure Security Group."
- 8. The security group determines who can connect to your instances. The default is "SSH" and 0.0.0.0/0. This means that any computer can connect to your instance over ssh. In practice, you would assign a restricted list of IP addresses that can access your cluster. However, for ease, you should leave it open. You'll also need to change the type to "All Traffic" so that we can access the instances over HTTP as well as SSH. Then, click "Review and Launch" to launch your instances.

At the end of the process, you'll be prompted to create and download a private key. This key will allow you to connect to your instances with SSH. If you don't download this, or delete it somehow, you won't be able to connect to your cluster. If you lose it, don't panic, but you'll have to shut down the instances and start up new ones. Name it whatever you like, I just used "hadoop." Click "View Instances" to see your instance booting up in the EC2 dashboard. You'll want to write down the public hostname, called "Public DNS" on the instance panel.

Configure SSH

You're going to access the instance by logging in remotely with SSH. If you are using Linux or Mac OS X, this should be available in the terminal. If you're on Windows, I suggest installing Git Bash which comes with Git. We are going to make heavy use of shell commands, so if you are unfamiliar with it, please take our Linux Command Line Basics course. Going forward, I am going to assume that you have access to a terminal that can connect to the instances.

Firstly, you'll want to change the permissions on the private key file, or you'll get an error when you try connecting through ssh. To change permission on Linux and OS X, \$ sudo chmod 600 /path/to/key_file.pem, where key_file.pem is the private key you created earlier.

Test logging in to the instance. Here instance_hostname is the public hostname of your instance:

```
$ ssh -i /path/to/key_file.pem ubuntu@instance_hostname
```

Now copy the private key to the instance:

```
$ scp -i /path/to/key_file.pem /path/to/key_file.pem
ubuntu@instance_hostname:~/.ssh
```

Install Java dan Hadoop

Update and upgrade:

```
$ sudo apt-get update && sudo apt-get dist-upgrade
Install OpenJDK:
```

```
$ sudo apt-get install openjdk-7-jdk
```

Download Hadoop from one of <u>these mirrors</u>. Change the version number appropriately:

```
$ wget http://apache.mirrors.tds.net/hadoop/common/hadoop-2.7.2/hadoop-
2.7.2.tar.gz -P ~/Downloads
```

Extract it to /usr/local:

```
$ sudo tar zxvf ~/Downloads/hadoop-* -C /usr/local
```

Rename the directory to just hadoop:

```
$ sudo mv /usr/local/hadoop-* /usr/local/hadoop
```

Set environment variables

Find Java with:

```
readlink -f $(which java)
```

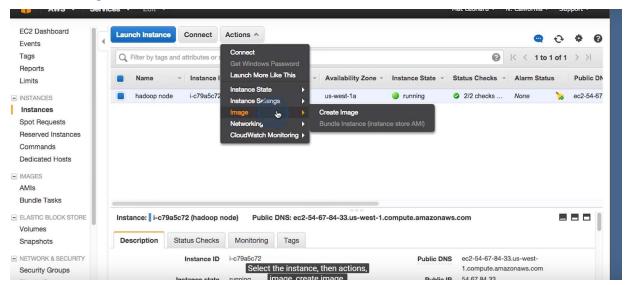
Environment variables:

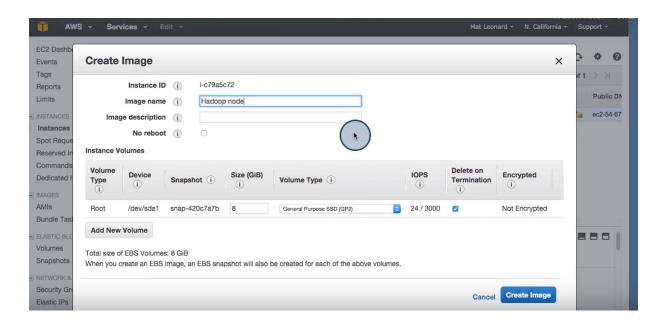
```
export JAVA_HOME=/path/to/java
export PATH=$PATH:$JAVA_HOME/bin
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export HADOOP_CONF_DIR=/usr/local/hadoop/etc/hadoop
```

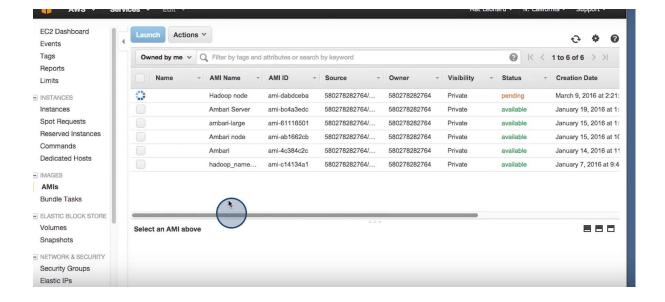
Load variables:

\$ source ~/.bashrc

Save to image







Launching nodes from the image

Now that you have an image with Java and Hadoop installed, you can use this as nodes in our cluster. Select the Hadoop node image, then click launch. Choose 4 instances, one will be the NameNode, the other three will be DataNodes.

Continue with configuration as before, remember to set the security to "All Traffic." When launching, select "Choose an existing key pair" and the private key that you created before. The new instances will use this private key for connecting over SSH.

You should see all four nodes in the instances panel. You'll need the public hostname of each instance often, so it's best to write each one in a file for reference.

To make logging into the instances more convenient, let's create an SSH config file. First, on your computer, create the file with \$ touch ~/.ssh/config, then edit ~/.ssh/config in your favorite editor. Add these lines to the file.

```
Host namenode

HostName namenode_public_hostname

User ubuntu

IdentityFile ~/.ssh/key_file.pem

Host datanode1

HostName datanode1_public_hostname

User ubuntu
```

```
IdentityFile ~/.ssh/key_file.pem

Host datanode2
HostName datanode2_public_hostname
User ubuntu
IdentityFile ~/.ssh/key_file.pem

Host datanode3
HostName datanode3_public_hostname
User ubuntu
IdentityFile ~/.ssh/key_file.pem
```

This file lets SSH associate a shorthand name with a hostname, a user, and the private key, so you don't have to type those in each time. This is assuming your private key key_file.pem is in .ssh. If it isn't be sure to move or copy it there: cp key_file ~/.ssh/key_file.pem. Now you can log into the NameNode with just \$ ssh namenode. Also, copy the config file to the NameNode:

```
$ scp ~/.ssh/config namenode:~/.ssh
```

You need to make sure the NameNode can connect to each DataNode over ssh without needing a password. You'll do this by creating a public key for the NameNode and adding it to each DataNode.

On the NameNode (remember \$ ssh namenode), create a public key and copy it into authorized_keys:

```
$ ssh-keygen -f ~/.ssh/id_rsa -t rsa -P ""
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

Copy authorized_keys to each DataNode to enable passwordless login. On the NameNode:

```
$ ssh datanode1 'cat >> ~/.ssh/authorized_keys' < ~/.ssh/id_rsa.pub
$ ssh datanode2 'cat >> ~/.ssh/authorized_keys' < ~/.ssh/id_rsa.pub
$ ssh datanode3 'cat >> ~/.ssh/authorized_keys' < ~/.ssh/id_rsa.pub</pre>
```

Test this by logging into a DataNode from the NameNode:

```
$ ssh datanode1
```

You can get back to the NameNode with

```
$ exit
```

- Configuring the cluster
- Hadoop environment variables
- Cluster-wide configuration
- NameNode Specific Configuration
- DataNode specific configuration

Configuring the cluster

You'll need the public hostname of each instance going forward, so it's really convenient to copy them down in a file. You can see each public hostname by selecting all the instances.

Hadoop has an intimidating number of configuration options. I'm only going to go through a few here, check the <u>Hadoop documentation</u> and <u>Hadoop: The Definitive</u> <u>Guide</u> for more information on configuration options.

Hadoop environment variables

The first file to take note of is hadoop-env.sh, it holds environment variables used by Hadoop. Here, you'll just need to set JAVA_HOME

```
export JAVA_HOME=/path/to/java
```

Cluster-wide configuration

First, you'll deal with the configuration on each node, then get into specific configurations for the NameNode and DataNodes. On each node, go to the Hadoop configuration folder, you should be able to get there with \$ cd \$HADOOP_CONF_DIR since we set that in .bashrc earlier. When editing these configuration files, you'll need root access so remember to use \$ sudo. In the configuration folder, edit core-site.xml:

```
<value>hdfs://namenode_public_hostname:9000</value>
  </property>
</configuration>
```

where namenode_public_hostname is the public hostname of your NameNode. This configuration fs.defaultFS tells the cluster nodes which machine the NameNode is on and that it will communicate on port 9000.

On each node, in yarn-site.xml you set options related to YARN, the resource manager:

Similarly with fs.defaultFS, yarn.resourcemanager.hostname sets the machine that the resource manager runs on.

On each node, copy mapred-site.xml from mapred-site.xml.template

```
$ sudo cp mapred-site.xml.template mapred-site.xml
and add:
```

```
</configuration>
```

Again, mapreduce.jobtracker.address sets the machine the job tracker runs on, and the port it communicates with. The other option here mapreduce.framework.name sets MapReduce to run on YARN.

NameNode specific configuration

Now, NameNode specific configuration, these will all be configured only on the NameNode. First, add the DataNode hostnames to /etc/hosts. You can get the hostname for each DataNode by entering hostname, or \$ echo (hostname) on each DataNode.

Now edit /etc/hosts and include these lines:

```
127.0.0.1 localhost

namenode_public_hostname namenode_hostname

datanode1_public_hostname datanode1_hostname

datanode2_public_hostname datanode2_hostname

datanode3_public_hostname datanode3_hostname
```

Now, on the NameNode, edit hdfs-site.xml:

dfs.replication sets how many times each data block is replicated across the cluster. dfs.namenode.name.dir sets the directory for storing NameNode data (.fsimage). You'll also have to create the directory to store the data:

```
$ sudo mkdir -p $HADOOP_HOME/data/hdfs/namenode
```

Next, you'll create the masters file in HADOOP_CONF_DIR. The masters file sets which machine the secondary namenode runs on. In your case, you'll have the secondary NameNode run on the same machine as the NameNode, so edit masters, add the

hostname of NameNode (**Note:** Not the public hostname, but the hostname you get from \$\\$\text{hostname}\$). Typically though, you would have the secondary NameNode run on a different machine than the primary NameNode.

Next, edit the slaves file in HADOOP_CONF_DIR, this file sets the machines that are DataNodes. In slaves, add the hostnames of each datanode (**Note:** Again, not the public hostname, but \$ hostname hostnames). The slaves file might already contain a line localhost, you should remove it, otherwise the NameNode would run as a DataNode too. It should look like this:

```
datanode1_hostname
datanode2_hostname
datanode3_hostname
```

Finally on the NameNode, change the owner of HADOOP_HOME to ubuntu:

```
$ sudo chown -R ubuntu $HADOOP_HOME
```

DataNode specific configuration

Now on **each** DataNode, edit HADOOP_CONF_DIR/hdfs-site.xml:

Again, this sets the directory where the data is stored on the DataNodes. And again, create the directory on each DataNode:

```
$ sudo mkdir -p $HADOOP_HOME/data/hdfs/datanode
and change the owner of the Hadoop directory
```

```
$ sudo chown -R ubuntu $HADOOP_HOME
```

Launch Hadoop Cluster

On the NameNode, format the file system, then start HDFS.

```
$ hdfs namenode -format
$ $HADOOP_HOME/sbin/start-dfs.sh
```

Start YARN:

```
$ $HADOOP_HOME/sbin/start-yarn.sh
```

Start the job history server:

```
$ $HADOOP_HOME/sbin/mr-jobhistory-daemon.sh start historyserver
```

To see the Java processes (Hadoop daemons for instance), enter

```
$ jps
```